

Indeed

August 3, 2020

1 Indeed Data Scraping Project

The goal of this project is to automate the Indeed Job searching process by allowing the user to input a city of their choice and returning a CSV file containing the cleaned job listings in the first 10 pages on Indeed. This will make the job searching process much easier for the user. The data scraped from the Indeed website can also be used to create an awesome dataframe for some data science project. An example would be to find and compare the salaries and ratings quartiles of data scientist positions in different cities.

1.1 Import Statements

```
[52]: from bs4 import BeautifulSoup as BSoup
import requests
import pandas as pd
```

1.2 Creating our Response and BeautifulSoup objects

Let's first start by attempting to create a dataframe from just one Indeed page URL. Note this is a sample URL.

```
[53]: URL = "https://www.indeed.com/q-Data-Scientist-l-San-Francisco,-CA-jobs.html?
        ↪vjk=bc7c0e642f6453f4"
request = requests.get(URL)
print(request)
```

<Response [200]>

Awesome! We got a response code 200, meaning that our request to the webpage was successful! Let's now view the HTML of the webpage and use BeautifulSoup to make it look nicer. I will comment out the prettify statement, which prints out a nice version of the HTML, so it won't display the html code because it is very, very long :) If you are interested in what it looks like, feel free to copy paste the URL into your browser, right click the webpage, and click Inspect.

```
[54]: page_html = BSoup(request.text, "html.parser")
# page_html.prettify()
```

Let's check how many job listings there are on this one webpage (not counting all the other pages for the listings). We will call each of these listings "containers" since they look like separate boxes when inspecting the page.

```
[55]: containers = page_html.findAll(name="div", attrs={"class": "row"})
len(containers)
```

[55]: 16

It looks like there are 15 jobs from this sample URL! ## Extracting data by looking at the HTML tags from the BeautifulSoup object Let's start out by extracting the job title.

```
[56]: # Scrapes the job titles from the BeautifulSoup object
def extract_job_title_from_result(soup):
    jobs = []
    for container in containers:
        for a in container.find_all(name="a", attrs={"data-tn-element": "jobTitle"}):
            jobs.append(a["title"])
    return jobs

extract_job_title_from_result(page_html)
```

```
[56]: ['Senior Applied Scientist',
      'Data Scientist',
      'Data Scientist',
      'Personalization Artificial Intelligence / Machine Learning (AI/ML) Data Solutions Consultant - Data Management Consultant 3',
      'Enterprise Architect (Data) (2020-9124)',
      'Associate Data Scientist I',
      'Data Scientist',
      'Data Scientist Entry Level - Pathrise Recruiting Partners',
      'Research Data Scientist',
      'Data Scientist: Data Visualization',
      'Junior Data Scientist - Pathrise Recruiting Partners',
      'Data Scientist',
      'Data Scientist',
      'Data Scientist',
      'Senior Data Scientist',
      'Data Scientist']
```

Let's do the same for the company.

```
[57]: # Scrapes the company from the BeautifulSoup object
def extract_company_from_result(soup):
    companies = []
    for container in containers:
        company = container.find_all(name="span", attrs={"class": "company"})
```

```

        if len(company) > 0:
            for b in company:
                companies.append(b.text.strip())
        else:
            test2 = div.find_all(name="span", attrs={"class": "↵
↵"result-link-source"})
            for span in test2:
                companies.append(span.text.strip())
        return(companies)

extract_company_from_result(page_html)

```

```

[57]: ['Spiketrapp',
      'Global Fishing Watch',
      'Blue Owl',
      'Wells Fargo Bank',
      'Fisher Investments',
      'Levi Strauss & Co.',
      'project AI',
      'Pathrise',
      'University of California San Francisco',
      'Kaiser Permanente',
      'Pathrise',
      'Applied Technology & Science (A-T-S)',
      'Common Networks',
      'Yelp',
      'MOLOCO',
      'Triplebyte']

```

Let's do the same for the salary.

```

[58]: # Scrapes the salary from the BeautifulSoup object
def extract_salary_from_result(soup):
    salaries = []
    for div in soup.find_all(name="div", attrs={"class": "row"}):
        div_two = div.find(name="span", attrs={'class': "salaryText"})
        if div_two == None:
            salaries.append("Not Available")
        else:
            salaries.append(div_two.text.strip())
    return salaries

extract_salary_from_result(page_html)

```

```

[58]: ['Not Available',
      '$45 - $65 an hour',
      '$250,000 - $375,000 a year',

```

```
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'$145,000 - $225,000 a year']
```

Finally, let's do the same for ratings.

```
[59]: # Scrapes the ratings from the BeautifulSoup object
def extract_ratings_from_result(soup):
    ratings = []
    for div in soup.find_all(name="div", attrs={"class": "row"}):
        div_two = div.find(name="span", attrs={'class': "ratingsContent"})
        if div_two == None:
            ratings.append("Not Available")
        else:
            ratings.append(div_two.text.strip())
    return ratings

extract_ratings_from_result(page_html)
```

```
[59]: ['Not Available',
'Not Available',
'Not Available',
'3.7',
'3.6',
'3.9',
'Not Available',
'Not Available',
'4.2',
'4.1',
'Not Available',
'Not Available',
'Not Available',
'3.5',
'Not Available',
'5.0']
```

Now let's build a dataframe by combining all the information we have so far!

```
[60]: example_df = pd.DataFrame(
    {"job_title": extract_job_title_from_result(page_html),
     "company": extract_company_from_result(page_html),
     "salary": extract_salary_from_result(page_html),
     "rating": extract_ratings_from_result(page_html)}
)
example_df
```

```
[60]:
```

	job_title \		company	salary \
0	Senior Applied Scientist		Spiketrapp	Not Available
1	Data Scientist		Global Fishing Watch	\$45 - \$65 an hour
2	Data Scientist		Blue Owl	\$250,000 - \$375,000 a year
3	Personalization Artificial Intelligence / Mach...		Wells Fargo Bank	Not Available
4	Enterprise Architect (Data) (2020-9124)		Fisher Investments	Not Available
5	Associate Data Scientist I		Levi Strauss & Co.	Not Available
6	Data Scientist		project AI	Not Available
7	Data Scientist Entry Level - Pathrise Recruti...		Pathrise	Not Available
8	Research Data Scientist		University of California San Francisco	Not Available
9	Data Scientist: Data Visualization		Kaiser Permanente	Not Available
10	Junior Data Scientist - Pathrise Recruiting Pa...		Pathrise	Not Available
11	Data Scientist		Applied Technology & Science (A-T-S)	Not Available
12	Data Scientist		Common Networks	Not Available
13	Data Scientist		Yelp	Not Available
14	Senior Data Scientist		MOLOCO	Not Available
15	Data Scientist		Triplebyte	\$145,000 - \$225,000 a year

	rating
0	Not Available
1	Not Available

2	Not Available
3	3.7
4	3.6
5	3.9
6	Not Available
7	Not Available
8	4.2
9	4.1
10	Not Available
11	Not Available
12	Not Available
13	3.5
14	Not Available
15	5.0

Awesome! It looks good except that we need to clean the salary series since it is not consistent with units (years and hour) and fix the style of the salary text. We won't worry too much about that right now. Let's now try to get all listings from the first 10 pages of the Indeed searches

1.3 Looping through the first 10 webpages of a specific city.

```
[61]: # Limit to first 10 pages counting from 0 in increments of 10 to match the URL
      ↪pattern
      limit = 100

      # Headers for the data we want to extract
      columns = ["job_title", "company", "salary", "rating"]
```

```
[62]: # We want the user to be able to input a city of their choice, so we will test
      ↪it using the input method
      a = input()
      city_selection = [a]
```

San Francisco

```
[63]: # Creates a dataframe
      sample_df = pd.DataFrame(columns=columns)

      # Loops through the 10 webpages for the selected city and gets the information
      ↪similar to the functions we created above
      for city in city_selection:
          for start in range(0, limit, 10):
              page = requests.get("http://www.indeed.com/jobs?
              ↪q=data+scientist+%2420%2C000&l=" + str(city) + "&start=" + str(start))
              soup = BSoup(page.text, "lxml", from_encoding="utf-8")
```

```

for div in soup.find_all(name="div", attrs={"class": "row"}):
    num = (len(sample_df) + 1)
    job_post = []

    for a in div.find_all(name="a", attrs={"data-tn-element": "
↪"jobTitle"}):
        job_post.append(a["title"])
        company = div.find_all(name="span", attrs={"class": "company"})
        if len(company) > 0:
            for b in company:
                job_post.append(b.text.strip())
        else:
            test2 = div.find_all(name="span", attrs={"class": "
↪"result-link-source"})
            for span in test2:
                job_post.append(span.text)

    div_two = div.find(name="span", attrs={"class": "salaryText"})
    if div_two == None:
        job_post.append("Not Available")
    else:
        job_post.append(div_two.text.strip())

    div_three = div.find(name="span", attrs={"class": "ratingsContent"})
    if div_three == None:
        job_post.append("Not Available")
    else:
        job_post.append(div_three.text.strip())

    sample_df.loc[num] = job_post

```

Awesome, now that our dataframe is created, let's check it out! We will look at the first and last 5 rows to avoid displaying all of it.

Note: Some of the rows in the notebook file uploaded on GitHub may include some tags and information on GitHub, which is a problem of displaying the data on a new row (you can see this by the rows skipping). This problem is not seen on the notebook itself. If you would like to see the normal dataframe, check out the PDF file of the notebook or the CSV file of the data.

```
[64]: sample_df.head()
```

```

[64]:
1          Senior Applied Scientist      Spiketrapp
2          Data Scientist               Triplebyte
3          Data Scientist               Blue Owl
4          Machine Learning Engineer    Triplebyte

```

5 Data Solutions Sr. Consultant/Personalization ... Wells Fargo Bank

	salary	rating
1	Not Available	Not Available
2	\$145,000 - \$225,000 a year	5.0
3	\$250,000 - \$375,000 a year	Not Available
4	\$150,000 - \$250,000 a year	5.0
5	Not Available	3.7

```
[65]: sample_df.tail()
```

```
[65]:
```

	job_title	company \
149	Senior Data Analyst	Good Eggs
150	Data and Evaluation Contractor	AI4ALL
151	Staff Data Scientist - Activision Blizzard Media	King.com
152	Senior Data Scientist & Modeler	SoFi
153	Data Scientist - Global Business Operations	Splunk

	salary	rating
149	Not Available	3.0
150	Not Available	Not Available
151	Not Available	Not Available
152	Not Available	3.2
153	Not Available	4.2

1.4 Data Cleaning

Let's first fix the salary series since there is a bolding issue with the Markdown concatenation of the '\$', the first number of the range, and '-'.

```
[66]: lst = []

for i in range(1, len(sample_df['salary']) + 1):
    # 6 figure salary a year with range
    if len(sample_df['salary'][i]) == 26:
        lst.append(sample_df['salary'][i][0:8] + ' - ' +
↪sample_df['salary'][i][12:19] + ' a year')

    # 5 figure salary a year with range
    elif len(sample_df['salary'][i]) == 25:
        lst.append(sample_df['salary'][i][0:7] + ' - ' +
↪sample_df['salary'][i][11:18] + ' a year')

    # 6 figure salary a year no range
    elif len(sample_df['salary'][i]) == 15:
        lst.append(sample_df['salary'][i][0:8] + ' a year')
```



```

# 2 figure salary a hour no range
elif len(sample_df['salary'][i]) == 17:
    lst.append(sample_df['salary'][i][0:3] + ' - ' +
→sample_df['salary'][i][7:9] + ' an hour')

# a 3 salary figure a hour no range
elif len(sample_df['salary'][i]) == 18:
    lst.append(sample_df['salary'][i][0:3] + ' - ' +
→sample_df['salary'][i][7:9] + ' an hour')

# 4 figure salary a month no range
elif len(sample_df['salary'][i]) == 14:
    lst.append(sample_df['salary'][i][0:6] + ' a month')

else:
    lst.append(sample_df['salary'][i])

sample_df['salary'] = lst

```

```
[67]: sample_df.head()
```

```
[67]:
```

	job_title	company \
1	Senior Applied Scientist	Spiketrap
2	Data Scientist	Triplebyte
3	Data Scientist	Blue Owl
4	Machine Learning Engineer	Triplebyte
5	Data Solutions Sr. Consultant/Personalization ...	Wells Fargo Bank

	salary	rating
1	Not Available	Not Available
2	\$145,000 - 225,000 a year	5.0
3	\$250,000 - 375,000 a year	Not Available
4	\$150,000 - 250,000 a year	5.0
5	Not Available	3.7

```
[68]: sample_df.tail()
```

```
[68]:
```

	job_title	company \
149	Senior Data Analyst	Good Eggs
150	Data and Evaluation Contractor	AI4ALL
151	Staff Data Scientist - Activision Blizzard Media	King.com
152	Senior Data Scientist & Modeler	SoFi
153	Data Scientist - Global Business Operations	Splunk

	salary	rating
149	Not Available	3.0

```

150 Not Available Not Available
151 Not Available Not Available
152 Not Available 3.2
153 Not Available 4.2

```

Next, we are going to clean this data and then convert the dataframe into a CSV file. Let's start by cleaning the salary series to the correct rates. We will convert them into dollars a year. There are a lot of different cases involving different units and ranges, and some cases may not be covered. I will focus on covering cases displayed in San Francisco. We will also make the \$ sign consistent among rows and different units.

Note: the string manipulation below assumes that hourly salaries are two digits and monthly to be in the thousands since salaries for these jobs. We can safely make this assumption for now as annual income for this position is usually 50k-200k.

```

[69]: result = sample_df['salary']

for index, item in enumerate(sample_df["salary"]):
    # Ranges in hour -> convert to year assuming 8 hrs a day, 5 times a week
    if "hour" in item and '-' in item:
        lower = int(item[1:3])*8*365
        upper = int(item[6:8])*8*365
        result[index + 1] = "$" + "{:,}".format(lower) + " - " + "{:,}".
        ↪format(upper) + " a year"

    # No range in hour -> convert to year assuming 8 hrs a day, 5 times a week
    elif "hour" in item and '-' not in item:
        salary = int(item[1:3])*8*365
        result[index + 1] = "$" + "{:,}".format(salary) + " a year"

    # No range in month -> convert to year assuming 8 hrs a day, 5 times a week
    elif 'month' in item:
        no_range = int(item[1:2] + item[3:6])*12
        result[index + 1] = "$" + "{:,}".format(no_range) + " a year"

    # Already in year
    else:
        result[index + 1] = item

result

```

```

[69]: 1          Not Available
2    $145,000 - 225,000 a year
3    $250,000 - 375,000 a year
4    $150,000 - 250,000 a year
5          Not Available
6          Not Available

```

7	Not Available
8	Not Available
9	Not Available
10	Not Available
11	Not Available
12	Not Available
13	Not Available
14	Not Available
15	Not Available
16	\$131,400 - 189,800 a year
17	\$145,000 - 225,000 a year
18	\$131,400 - 189,800 a year
19	\$60,000 a year
20	Not Available
21	Not Available
22	\$120,000 a year
23	Not Available
24	Not Available
25	Not Available
26	Not Available
27	Not Available
28	Not Available
29	Not Available
30	\$250,000 - 375,000 a year
	...
124	Not Available
125	Not Available
126	Not Available
127	Not Available
128	Not Available
129	Not Available
130	Not Available
131	Not Available
132	Not Available
133	Not Available
134	\$150,000 - 250,000 a year
135	\$60,000 a year
136	\$145,000 - 225,000 a year
137	\$131,400 - 189,800 a year
138	\$250,000 - 375,000 a year
139	Not Available
140	Not Available
141	Not Available
142	\$132,476 - 163,000 a year
143	Not Available
144	Not Available
145	Not Available

```

146          Not Available
147          Not Available
148          Not Available
149          Not Available
150          Not Available
151          Not Available
152          Not Available
153          Not Available
Name: salary, Length: 153, dtype: object

```

Let's check our dataframe now for our conversion rates, which should all be \$ a year.

```
[73]: sample_df
```

```

[73]:                                     job_title \
1                Senior Applied Scientist
2                Data Scientist
3                Data Scientist
4                Machine Learning Engineer
5  Data Solutions Sr. Consultant/Personalization ...
6                Associate Data Scientist I
7                Data Scientist
8                Research Data Scientist
9  Data Scientist Entry Level - Pathrise Recruti...
10               Data Scientist: Data Visualization
11  Junior Data Scientist - Pathrise Recruiting Pa...
12               Data Scientist
13  Data Scientist / Quantitative Research
14               Data Scientist
15               Data Scientist
16               Data Scientist
17               Data Scientist
18               Data Scientist
19  Financial Analyst Summer Intern - Ideal for a ...
20  Enterprise Architect (Data) (2020-9124)
21               Senior Data Scientist
22               Data Scientist
23  Data Scientist, Marketing
24  Data Scientist, Legal Policy & Economics
25  Data Scientist, Machine Learning innovator
26  Data Scientist - Shop Recommendations
27  Data Scientist - Experimentation (Contract Pos...
28               Data Scientist
29  Data Science Intern - Remote
30               Data Scientist
..               ...
124  Machine Learning Engineer Entry Level - Pathri...

```

125 Computer Vision Data Scientist
 126 Senior Data Scientist
 127 Data Scientist - Regional Merchandising Strategy
 128 Data Scientist/Research Engineer
 129 Data Science Instructional Designer and Analyst
 130 Senior Data Analyst
 131 Sr. Pharmacy Data Analyst
 132 Data Scientist, Medidata - Core
 133 Sr. Data Analyst, Twitter Service Tech
 134 Machine Learning Engineer
 135 Financial Analyst Summer Intern - Ideal for a ...
 136 Data Scientist
 137 Data Scientist
 138 Data Scientist
 139 Principal Data Scientist
 140 Data Engineer
 141 Investment Banking Analyst Intern, Summer 2021
 142 Machine Learning Researcher - Multi View & Seg...
 143 Marketing and Business Analytics Interns
 144 Senior Data Scientist - Marketing
 145 Data Engineer - Python Programmer
 146 Staff Data Scientist - Global Payments & Fraud...
 147 Machine Learning Engineers (Multiple Opportuni...
 148 Data Science Manager- Health
 149 Senior Data Analyst
 150 Data and Evaluation Contractor
 151 Staff Data Scientist - Activision Blizzard Media
 152 Senior Data Scientist & Modeler
 153 Data Scientist - Global Business Operations

company \
 1 Spiketrap
 2 Triplebyte
 3 Blue Owl
 4 Triplebyte
 5 Wells Fargo Bank
 6 Levi Strauss & Co.
 7 project AI
 8 University of California San Francisco
 9 Pathrise
 10 Kaiser Permanente
 11 Pathrise
 12 Applied Technology & Science (A-T-S)
 13 PicnicHealth
 14 Yelp
 15 Common Networks
 16 Global Fishing Watch

17	Triplebyte
18	Global Fishing Watch
19	MPL Brands
20	Fisher Investments
21	MOLOCO
22	GradTests (gradtests.com)
23	Twitch
24	Uber
25	Standard Chartered
26	Stitch Fix
27	Getty Images
28	Deep Labs
29	Interview Query
30	Blue Owl
..	...
124	Pathrise
125	Enlitic
126	Landing
127	Wish
128	Apixio
129	University of California San Francisco
130	Life360
131	University of California San Francisco Medical...
132	Medidata Solutions
133	Twitter
134	Triplebyte
135	MPL Brands
136	Triplebyte
137	Global Fishing Watch
138	Blue Owl
139	Tapjoy
140	CommonStock
141	GCA
142	Fyusion
143	UpCounsel LLC
144	Opendoor
145	Mondo
146	Sony Interactive Entertainment PlayStation
147	Grelock Partners
148	Twitter
149	Good Eggs
150	AI4ALL
151	King.com
152	SoFi
153	Splunk

salary	rating
--------	--------

1		Not Available	Not Available
2	\$145,000 - 225,000 a year		5.0
3	\$250,000 - 375,000 a year	Not Available	
4	\$150,000 - 250,000 a year		5.0
5		Not Available	3.7
6		Not Available	3.9
7		Not Available	Not Available
8		Not Available	4.2
9		Not Available	Not Available
10		Not Available	4.1
11		Not Available	Not Available
12		Not Available	Not Available
13		Not Available	Not Available
14		Not Available	3.5
15		Not Available	Not Available
16	\$131,400 - 189,800 a year	Not Available	
17	\$145,000 - 225,000 a year		5.0
18	\$131,400 - 189,800 a year	Not Available	
19	\$60,000 a year	Not Available	
20		Not Available	3.6
21		Not Available	Not Available
22	\$120,000 a year	Not Available	
23		Not Available	4.4
24		Not Available	3.7
25		Not Available	4.1
26		Not Available	3.2
27		Not Available	3.9
28		Not Available	3.7
29		Not Available	Not Available
30	\$250,000 - 375,000 a year	Not Available	
..	
124		Not Available	Not Available
125		Not Available	Not Available
126		Not Available	3.8
127		Not Available	3.8
128		Not Available	4.3
129		Not Available	4.2
130		Not Available	4.5
131		Not Available	4.2
132		Not Available	3.7
133		Not Available	4.1
134	\$150,000 - 250,000 a year		5.0
135	\$60,000 a year	Not Available	
136	\$145,000 - 225,000 a year		5.0
137	\$131,400 - 189,800 a year	Not Available	
138	\$250,000 - 375,000 a year	Not Available	
139		Not Available	3.3

140		Not Available	Not Available
141		Not Available	3.1
142	\$132,476 - 163,000 a year	Not Available	
143		Not Available	Not Available
144		Not Available	3.0
145		Not Available	Not Available
146		Not Available	3.7
147		Not Available	Not Available
148		Not Available	4.1
149		Not Available	3.0
150		Not Available	Not Available
151		Not Available	Not Available
152		Not Available	3.2
153		Not Available	4.2

[153 rows x 4 columns]

Awesome! It looks like it worked properly. For example, line 2 with company Global Fish Watch was converted from 45-65 to 131,400-189,800. Let's now finally convert this Pandas dataframe into a CSV file and check it out!

```
[74]: sample_df.to_csv('indeed.csv', index=False)
```

```
[ ]:
```