

# Indeed

June 3, 2021

## 1 Scraping Indeed Listings

The goal of this project is to scrape information off of a job listing on Indeed's website. This process would include the user inputting a city of their choice and returning a CSV file containing information of data science job listings in the first 10 pages of the search result. This will make the job searching process much easier for the user and the data scraped from the can also be used to create an awesome dataframe for some data science project. An example would be to find and compare the salaries and ratings quartiles of data scientist positions in different cities.

Through this project, I hope to gain exposure to web scraping through the use of BeautifulSoup.

### 1.1 Import Statements

```
[1]: from bs4 import BeautifulSoup as BSoup
import requests
import pandas as pd
```

### 1.2 Creating our Response and BeautifulSoup objects

Let's first start by attempting to create a dataframe from just one Indeed page URL. Note this is a sample URL.

```
[2]: URL = "https://www.indeed.com/jobs?q=Data+Scientist&l=New+York"
request = requests.get(URL)
print(request)
```

<Response [200]>

Awesome! We got a response code 200, meaning that our request to the webpage was successful! Let's now view the HTML of the webpage and use BeautifulSoup to make it look nicer. I will comment out the prettify statement, which prints out a nice version of the HTML, so it won't display the html code because it is very, very long :) If you are interested in what it looks like, feel free to copy paste the URL into your browser, right click the webpage, and click Inspect.

```
[3]: page_html = BSoup(request.text, "html.parser")
# page_html.prettify()
```

Let's check how many job listings there are on this specific webpage (not counting all the other pages for the listings). We will call these each of these listings "containers" since they look like separate boxes when inspecting the page.

```
[5]: containers = page_html.findAll(name="div", attrs={"class": "row"})
len(containers)
```

```
[5]: 15
```

It looks like there are 15 job from this sample URL! ## Extracting data by looking at the HTML tags from the BeautifulSoup object Let's start out by extracting the job title.

```
[6]: # Scrapes the job titles from the BeautifulSoup object
def extract_job_title_from_result(soup):
    jobs = []
    for div in soup.find_all(name="div", attrs={"class": "row"}):
        for a in div.find_all(name="a", attrs={"data-tn-element": "jobTitle"}):
            jobs.append(a["title"])
    return(jobs)

extract_job_title_from_result(page_html)
```

```
[6]: ['Data Scientist',
      'Data Scientist',
      'Data Scientist',
      'Data Scientist',
      'Data Scientist, Podcasts',
      'Senior Data Scientist',
      'Data Scientist',
      'junior Data Scientist',
      'Data Scientist',
      'Data Scientist',
      'Junior Data Scientist',
      'Data Scientist',
      'Data Scientist, SEO Analytics',
      'Data Scientist',
      'Data Scientist']
```

Let's do the same for the company.

```
[7]: # Scrapes the company from the BeautifulSoup object
def extract_company_from_result(soup):
    companies = []
    for container in containers:
        company = container.find_all(name="span", attrs={"class": "company"})
        if len(company) > 0:
            for b in company:
                companies.append(b.text.strip())
```

```

        else:
            test2 = div.find_all(name="span", attrs={"class": "
→"result-link-source"})
            for span in test2:
                companies.append(span.text.strip())
            return(companies)

extract_company_from_result(page_html)

```

```

[7]: ['PrimeNeuro',
      'Digital Republic Talent',
      'Robert Half',
      'Hinge',
      'Spotify',
      'Corning',
      'Carta',
      'CFSB',
      'Betterview',
      'Pluto TV',
      'Decode_M',
      'Boll & Branch',
      'Reddit',
      'Source Enterprises',
      'Gannett']

```

Let's do the same for the salary.

```

[8]: # Scrapes the salary from the BeautifulSoup object
def extract_salary_from_result(soup):
    salaries = []
    for div in soup.find_all(name="div", attrs={"class": "row"}):
        div_two = div.find(name="span", attrs={'class': "salaryText"})
        if div_two == None:
            salaries.append("Not Available")
        else:
            salaries.append(div_two.text.strip())
    return salaries

extract_salary_from_result(page_html)

```

```

[8]: ['$95,000 - $105,000 a year',
      '$100,000 - $140,000 a year',
      '$170,000 - $175,000 a year',
      'Not Available',
      'Not Available',
      'Not Available',
      'Not Available',

```

```
'$80,000 - $90,000 a year',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'Not Available']
```

Finally, let's do the same for ratings.

```
[9]: # Scrapes the ratings from the BeautifulSoup object
def extract_ratings_from_result(soup):
    ratings = []
    for div in soup.find_all(name="div", attrs={"class": "row"}):
        div_two = div.find(name="span", attrs={'class': "ratingsContent"})
        if div_two == None:
            ratings.append("Not Available")
        else:
            ratings.append(div_two.text.strip())
    return ratings

extract_ratings_from_result(page_html)
```

```
[9]: ['Not Available',
'Not Available',
'3.9',
'4.8',
'4.3',
'3.8',
'3.8',
'4.4',
'Not Available',
'Not Available',
'Not Available',
'Not Available',
'4.0',
'4.3',
'3.0']
```

Now let's build a dataframe by combining all the information we have so far!

```
[10]: example_df = pd.DataFrame(
    {"job_title": extract_job_title_from_result(page_html),
    "company": extract_company_from_result(page_html),
    "salary": extract_salary_from_result(page_html),
    "rating": extract_ratings_from_result(page_html)}
```

```
)
example_df
```

```
[10]:
```

	job_title	company \
0	Data Scientist	PrimeNeuro
1	Data Scientist	Digital Republic Talent
2	Data Scientist	Robert Half
3	Data Scientist	Hinge
4	Data Scientist, Podcasts	Spotify
5	Senior Data Scientist	Corning
6	Data Scientist	Carta
7	junior Data Scientist	CFSB
8	Data Scientist	Betterview
9	Data Scientist	Pluto TV
10	Junior Data Scientist	Decode_M
11	Data Scientist	Boll & Branch
12	Data Scientist, SEO Analytics	Reddit
13	Data Scientist	Source Enterprises
14	Data Scientist	Gannett

  

	salary	rating
0	\$95,000 - \$105,000 a year	Not Available
1	\$100,000 - \$140,000 a year	Not Available
2	\$170,000 - \$175,000 a year	3.9
3	Not Available	4.8
4	Not Available	4.3
5	Not Available	3.8
6	Not Available	3.8
7	\$80,000 - \$90,000 a year	4.4
8	Not Available	Not Available
9	Not Available	Not Available
10	Not Available	Not Available
11	Not Available	Not Available
12	Not Available	4.0
13	Not Available	4.3
14	Not Available	3.0

Awesome! It looks good except that we need to clean the salary series since it is not consistent with units (years and hour) and fix the style of the salary text. We won't worry too much about that right now. Let's now try to get all listings from the first 10 pages of the Indeed searches

### 1.3 Looping through the first 5 webpages of a specific city.

```
[11]: # Limit to first 10 pages counting from 0 in increments of 10 to match the URL
      ↪pattern
      limit = 50
```

```
# Headers for the data we want to extract
columns = ["job_title", "company", "salary", "rating"]
```

```
[14]: # We want the user to be able to input a city of their choice, so we will test
      ↪it using the input method
      a = input()
      city_selection = [a]
```

San Francisco

```
[15]: # Creates a dataframe
      sample_df = pd.DataFrame(columns=columns)

      # Loops through the 10 webpages for the selected city and gets the information
      ↪similar to the functions we created above
      for city in city_selection:
          for start in range(0, limit, 10):
              page = requests.get("https://www.indeed.com/jobs?q=Data+Scientist&l=" +
              ↪str(city) + "&start=" + str(start))
              soup = BSoup(page.text, "lxml", from_encoding="utf-8")

              for div in soup.find_all(name="div", attrs={"class": "row"}):
                  num = (len(sample_df) + 1)
                  job_post = []

                  for a in div.find_all(name="a", attrs={"data-tn-element":
                  ↪"jobTitle"}):
                      job_post.append(a["title"])
                  company = div.find_all(name="span", attrs={"class": "company"})
                  if len(company) > 0:
                      for b in company:
                          job_post.append(b.text.strip())
                  else:
                      test2 = div.find_all(name="span", attrs={"class":
                      ↪"result-link-source"})
                      for span in test2:
                          job_post.append(span.text)

                  div_two = div.find(name="span", attrs={"class": "salaryText"})
                  if div_two == None:
```

```

        job_post.append("Not Available")
    else:
        job_post.append(div_two.text.strip())

    div_three = div.find(name="span", attrs={"class": "ratingsContent"})
    if div_three == None:
        job_post.append("Not Available")
    else:
        job_post.append(div_three.text.strip())

    sample_df.loc[num] = job_post

```

Awesome, now that our dataframe is created, let's check it out! We will look at the first and last 5 rows to avoid displaying all of it.

**Note:** Some of the rows in the notebook file uploaded on GitHub may include some tags and information on GitHub, which is a problem of displaying the data on a new row (you can see this by the rows skipping). This problem is not seen on the notebook itself. If you would like to see the normal dataframe, check out the PDF file of the notebook or the CSV file of the data.

```
[16]: sample_df.head()
```

```

[16]:
      job_title      company \
1      Staff Data Scientist      Metromile
2  Principal Data Scientist - Candidate Recommend...      Indeed
3      Data Scientist III  Thermo Fisher - America
4      Staff Data Scientist  Thermo Fisher - America
5      Data Scientist      PrimeNeuro

      salary      rating
1      Not Available      2.8
2  $187,000 - $231,000 a year      4.3
3      Not Available      3.5
4      Not Available      3.5
5  $95,000 - $105,000 a year  Not Available

```

```
[17]: sample_df.tail()
```

```

[17]:
      job_title      company \
45      Data Scientist  Zipcar, Inc.
46      Data Scientist  Commonstock
47  Data Scientist, Customer Intelligence - Opport...      VMware
48      Data Scientist      RAPP
49  Data Analyst - Disqus  Zeta Global

      salary      rating
45  Not Available      3.5

```

```

46 Not Available Not Available
47 Not Available 4.0
48 Not Available 3.5
49 Not Available 2.6

```

## 1.4 Data Cleaning

Let's first fix the salary series since there is a bolding issue with the Markdown concatenation of the '\$', the first number of the range, and '-'.

```

[18]: lst = []

for i in range(1, len(sample_df['salary']) + 1):
    # 6 figure salary a year with range
    if len(sample_df['salary'][i]) == 26:
        lst.append(sample_df['salary'][i][0:8] + ' - ' +
→sample_df['salary'][i][12:19] + ' a year')

    # 5 figure salary a year with range
    elif len(sample_df['salary'][i]) == 25:
        lst.append(sample_df['salary'][i][0:7] + ' - ' +
→sample_df['salary'][i][11:18] + ' a year')

    # 6 figure salary a year no range
    elif len(sample_df['salary'][i]) == 15:
        lst.append(sample_df['salary'][i][0:8] + ' a year')

    # 2 figure salary a hour no range
    elif len(sample_df['salary'][i]) == 17:
        lst.append(sample_df['salary'][i][0:3] + ' - ' +
→sample_df['salary'][i][7:9] + ' an hour')

    # a 3 salary figure a hour no range
    elif len(sample_df['salary'][i]) == 18:
        lst.append(sample_df['salary'][i][0:3] + ' - ' +
→sample_df['salary'][i][7:9] + ' an hour')

    # 4 figure salary a month no range
    elif len(sample_df['salary'][i]) == 14:
        lst.append(sample_df['salary'][i][0:6] + ' a month')

    else:
        lst.append(sample_df['salary'][i])

sample_df['salary'] = lst

```



```
[19]: sample_df.head()
```

```
[19]:
```

	job_title	company \
1	Staff Data Scientist	Metromile
2	Principal Data Scientist - Candidate Recommend...	Indeed
3	Data Scientist III	Thermo Fisher - America
4	Staff Data Scientist	Thermo Fisher - America
5	Data Scientist	PrimeNeuro

  

	salary	rating
1	Not Available	2.8
2	\$187,000 - 231,000 a year	4.3
3	Not Available	3.5
4	Not Available	3.5
5	\$95,000 - 105,000 a year	Not Available

```
[20]: sample_df.tail()
```

```
[20]:
```

	job_title	company \
45	Data Scientist	Zipcar, Inc.
46	Data Scientist	Commonstock
47	Data Scientist, Customer Intelligence - Opport...	VMware
48	Data Scientist	RAPP
49	Data Analyst - Disqus	Zeta Global

  

	salary	rating
45	Not Available	3.5
46	Not Available	Not Available
47	Not Available	4.0
48	Not Available	3.5
49	Not Available	2.6

Next, we are going to clean this data and then convert the dataframe into a CSV file. Let's start by cleaning the salary series to the correct rates. We will convert them into dollars a year. There are a lot of different cases involving different units and ranges, and some cases may not be covered. I will focus on covering cases displayed in San Francisco. We will also make the \$ sign consistent among rows and different units.

**Note: the string manipulation below assumes that hourly salaries are two digits and monthly to be in the thousands since salaries for these jobs. We can safely make this assumption for now as annual income for this position is usually 50k-200k.**

```
[21]: result = sample_df['salary']

for index, item in enumerate(sample_df["salary"]):
    # Ranges in hour -> convert to year assuming 8 hrs a day, 5 times a week
    if "hour" in item and '-' in item:
```

```

        lower = int(item[1:3])*8*365
        upper = int(item[6:8])*8*365
        result[index + 1] = "$" + "{:,}".format(lower) + " - " + "{:,}".
→format(upper) + " a year"

# No range in hour -> convert to year assuming 8 hrs a day, 5 times a week
elif "hour" in item and '-' not in item:
    salary = int(item[1:3])*8*365
    result[index + 1] = "$" + "{:,}".format(salary) + " a year"

# No range in month -> convert to year assuming 8 hrs a day, 5 times a week
elif 'month' in item:
    no_range = int(item[1:2] + item[3:6])*12
    result[index + 1] = "$" + "{:,}".format(no_range) + " a year"

# Already in year
else:
    result[index + 1] = item

result

```

```

[21]: 1          Not Available
      2    $187,000 - 231,000 a year
      3          Not Available
      4          Not Available
      5    $95,000 - 105,000 a year
      6    $130,000 - 156,000 a year
      7          Not Available
      8          Not Available
      9          Not Available
     10          Not Available
     11          Not Available
     12          Not Available
     13    $85,000 - 125,000 a year
     14          Not Available
     15    $120,000 - 160,000 a year
     16          Not Available
     17    $95,000 - 105,000 a year
     18          Not Available
     19          Not Available
     20          Not Available
     21          Not Available
     22          Not Available
     23          Not Available
     24          Not Available
     25          Not Available
     26          Not Available

```

```

27          Not Available
28          Not Available
29          Not Available
30          Not Available
31          Not Available
32          Not Available
33          Not Available
34          Not Available
35          Not Available
36          Not Available
37          Not Available
38          Not Available
39          Not Available
40          Not Available
41          Not Available
42          Not Available
43          Not Available
44    $150,000 - 180,000 a year
45          Not Available
46          Not Available
47          Not Available
48          Not Available
49          Not Available
Name: salary, dtype: object

```

Let's check our dataframe now for our conversion rates, which should all be \$ a year.

```
[22]: sample_df
```

```

[22]:
           job_title \
1          Staff Data Scientist
2  Principal Data Scientist - Candidate Recommend...
3          Data Scientist III
4          Staff Data Scientist
5          Data Scientist
6  Senior Data Scientist - Moderation Engineering
7  Senior Data Scientist - Machine Learning
8          Data Scientist
9          Data Scientist
10         Data Scientist
11         Data Scientist
12    Data Scientist - Intermediate
13         Data Scientist
14         Data Scientist
15    Abl Schools | Data Scientist
16    Associate Data Scientist
17         Data Scientist

```

18 Artificial Intelligence/Machine Learning Data ...  
 19 Principal Data Scientist - Telecommute  
 20 Staff Data Scientist - Product  
 21 Machine Learning Data Scientist  
 22 Data Scientist, Machine Learning  
 23 Data Scientist  
 24 Data Scientist, Analytics - Messaging Graph  
 25 Data Scientist  
 26 Research Data Scientist  
 27 Data Product Associate  
 28 Staff Data Scientist  
 29 Vitria Data Scientist  
 30 Data Science Consultant  
 31 Data Scientist - Research & Economics  
 32 Senior Data Analyst-Marketing  
 33 Data Scientist (Digital Analytics & Monetization)  
 34 Data Scientist  
 35 Data Scientist, Analytics  
 36 Data Scientist  
 37 Data Scientist  
 38 Associate Software Development Engineer  
 39 Data Scientist  
 40 Data Management Associate  
 41 Data Scientist  
 42 Data Scientist  
 43 Data Scientist Director  
 44 Sr. Data Scientist/Machine Learning Engineer  
 45 Data Scientist  
 46 Data Scientist  
 47 Data Scientist, Customer Intelligence - Opport...  
 48 Data Scientist  
 49 Data Analyst - Disqus

	company	salary \
1	Metromile	Not Available
2	Indeed	\$187,000 - 231,000 a year
3	Thermo Fisher - America	Not Available
4	Thermo Fisher - America	Not Available
5	PrimeNeuro	\$95,000 - 105,000 a year
6	Indeed	\$130,000 - 156,000 a year
7	Blue Owl	Not Available
8	Gap Inc.	Not Available
9	Carta	Not Available
10	First Republic Bank	Not Available
11	VIZIO, Inc.	Not Available
12	Bayer	Not Available
13	The Beans	\$85,000 - 125,000 a year

14	Grammarly, Inc.	Not Available
15	Abl Schools	\$120,000 - 160,000 a year
16	Gap Inc.	Not Available
17	PrimeNeuro	\$95,000 - 105,000 a year
18	Mitre Corporation	Not Available
19	UnitedHealth Group	Not Available
20	Twitter	Not Available
21	Oura	Not Available
22	Carta	Not Available
23	Komodo Health	Not Available
24	Facebook	Not Available
25	Cognizant Technology Solutions	Not Available
26	University of California San Francisco	Not Available
27	Levi Strauss & Co.	Not Available
28	Opendoor	Not Available
29	Vitria Technology	Not Available
30	Accenture	Not Available
31	Uber	Not Available
32	Kaiser Permanente	Not Available
33	Sony Interactive Entertainment PlayStation	Not Available
34	Tubi	Not Available
35	First Place for Youth	Not Available
36	a.k.a. Brands	Not Available
37	Grid Dynamics	Not Available
38	NextEra Energy	Not Available
39	Aquabyte	Not Available
40	ClimateWorks Foundation	Not Available
41	Joby Aviation	Not Available
42	Second Genome	Not Available
43	Oracle	Not Available
44	Perfect Minds	\$150,000 - 180,000 a year
45	Zipcar, Inc.	Not Available
46	Commonstock	Not Available
47	VMware	Not Available
48	RAPP	Not Available
49	Zeta Global	Not Available

	rating
1	2.8
2	4.3
3	3.5
4	3.5
5	Not Available
6	4.3
7	Not Available
8	3.8
9	3.8

10	3.9
11	Not Available
12	4.2
13	Not Available
14	4.6
15	Not Available
16	3.8
17	Not Available
18	4.0
19	3.7
20	4.1
21	Not Available
22	3.8
23	3.3
24	4.1
25	3.9
26	4.2
27	3.9
28	4.2
29	3.3
30	4.0
31	3.7
32	4.1
33	3.7
34	Not Available
35	3.4
36	Not Available
37	Not Available
38	3.9
39	Not Available
40	Not Available
41	Not Available
42	4.5
43	3.8
44	Not Available
45	3.5
46	Not Available
47	4.0
48	3.5
49	2.6

Awesome! It looks like it works properly. Let's now finally convert this Pandas dataframe into a CSV file and check it out!

```
[23]: sample_df.to_csv('indeed.csv', index=False)
```

```
[ ]:
```